

Lecture 14

Cost Functions, Regularization and Weight Initialization

07 March 2016

Taylor B. Arnold
Yale Statistics
STAT 365/665

Yale

Notes:

- ▶ Problem set 5 is due a week from today

Outline for this week:

- ▶ Monday
 - ▶ walk through R reference implementation of back-propagation
 - ▶ address issues of slow learning
 - ▶ over-fitting / regularization
 - ▶ initializing weights
- ▶ Wednesday
 - ▶ quasi second order SGD
 - ▶ learning rate schedule
 - ▶ hyper-parameter selection heuristics

Today's Notes

For today, rather than re-writing someone else's notes, I am simply going to go through the notes on this page:

<http://neuralnetworksanddeeplearning.com/chap3.html>

The primary reason for this is that there are some great visualizations embedded on the page; also, as I have used the exact same notation as presented here, there should be no great confusion in doing so.

There are three summary slides here for reference purposes, which summarize the main points I am covering today.

Cross-entropy and soft-max

Due to the shape of the sigmoid neuron, weights that are very far from their optimal values learn slowly in a plain, vanilla network. Two ways to fix this are to use the cross-entropy cost-function, defined as:

$$C = - \sum_j [y_j \log(a_j^L) + (1 - y_j) \log(1 - a_j^L)]$$

For a single sample, and similarly for an entire mini-batch.

Another common approach is to define what is termed a softmax layer. The redefines the activations of the output layer, a^L , as follows:

$$a_j^L = \frac{e^{z_j^L}}{\sum_k e^{z_k^L}}$$

This has the additional benefit that the last layer is easily interpreted as a sequence of probabilities.

Regularization in Neural Networks

As the size of neural networks grow, the number of weights and biases can quickly become quite large. State of the art neural networks today often have billions of weight values. In order to avoid over-fitting, one common approach is to add a penalty term to the cost function. Common choices are the ℓ_2 -norm, given as:

$$C = C_0 + \lambda \sum_i w_i^2$$

Where C_0 is the unregularized cost, and the ℓ_1 -norm:

$$C = C_0 + \lambda \sum_i |w_i|.$$

The distinction between these is similar to the differences between lasso and ridge regression.

Dropout

A very different approach to avoiding over-fitting is to use an approach called *dropout*. Here, the output of a randomly chosen subset of the neurons are temporarily set to zero during the training of a given mini-batch. This makes it so that the neurons cannot overly adapt to the output from prior layers as these are not always present. It has enjoyed wide-spread adoption and massive empirical evidence as to its usefulness.

Initial weights

Rather than initializing the weights in the neural network by a standard normal Gaussian, it is better to initialize with Gaussian noise with zero mean but a standard deviation of $1/\sqrt{n}$ where n is the number of nodes from the previous layer that are incoming to this given node.

This has the effect of making the distribution of all of the weighted inputs, z_j^l , equal to standard Gaussian noise.